

Alejandro Moreno

AI / Backend Engineer

alex@alexlaguardia.dev (717) 376-7870 Pennsylvania, USA alexlaguardia.dev github.com/AlexlaGuardia
LinkedIn

Open-source work & products published as “Alex LaGuardia”

SUMMARY

AI and backend engineer who designs, ships, and distributes production systems end-to-end, specializing in agent governance and MCP infrastructure: authorization enforced outside the model, tamper-evident attribution of which human directed an agent, and red-team harnesses that prove the enforcement holds. Strongest in Python/FastAPI and LLM systems (routing, RAG, MCP, agentic loops) plus the backend plumbing that keeps them alive, with open-source tools on PyPI and the VS Code Marketplace, a multi-tenant SaaS, and merged pull requests in codebases I did not write.

SKILLS

AI / LLM: Multi-LLM routing (Groq, Claude, Gemini), MCP (servers + clients), agentic tool loops, RAG, embeddings/semantic memory

Backend: FastAPI, async workers, daemon & queue architecture, SQLite, REST/SSE, OAuth2/JWT, Stripe webhooks

Infra / DevOps: PM2 (multi-service), Cloudflare (tunnels, Pages, CF-for-SaaS), Linux, SSH hardening, multi-tenant isolation

AI Dev Tools: Claude Code (daily driver); ships most code through a self-built MCP tool layer

Languages: Python (expert), Rust (advanced), TypeScript, SQL, Lua

Frontend: Next.js, React, Tailwind CSS, App Router

Integrations: Stripe, Groq, OANDA, Alpaca, Meta Graph API, Resend, Cloudflare API

SELECTED PROJECTS

Agent-Governance Suite: Warden · Crumb · Siege

Python, FastAPI, MCP, RBAC, OAuth / JWT, Sigstore

warden.alexlaguardia.dev · crumb.alexlaguardia.dev · github.com/AlexlaGuardia/siege

- **Warden** governs: RBAC policy engine authorizes every agent tool call outside the model (region row-scoping, field redaction), so a compromised agent cannot exceed its grants; grounded eval harness and a live enforced-policy demo
- **Crumb** attributes: tamper-evident proof of which human directed an agent's tool call, verifiable across separate identity providers by fetching each issuer's key live (JWKS), never from the log under audit
- **Siege** proves it holds: runtime red-team harness that attacks a live MCP server as real roles, surfacing authz / RBAC bypasses a static manifest scan cannot see

Vigil: Observability & Awareness Infrastructure for AI Agents

Python, FastAPI, MCP, SQLite, Stripe

PyPI: vigil-agent · app.vigil-agent.com · github.com/AlexlaGuardia/Vigil

- Open-source package + hosted multi-tenant SaaS: one-line MCP instrumentation (latency p50/95/99, per-tool error rates, silent-failure detection) plus an awareness layer (signal bus, 90s daemon-compiled context, frame-based tool filtering, session handoff)
- Hosted tier: per-tenant SQLite isolation with LRU-cached connections, GitHub OAuth, hashed API keys, Stripe billing, ~8,400 lines, 311 tests, 4 PyPI releases

Critik: AI Code Security Scanner

Python, Tree-sitter, Groq / Llama 3.3 70B

PyPI: kritik · VS Code Marketplace · GitHub Action · kritik.dev

- Two-pass scanner: regex + AST flag candidates, then an LLM reviews each finding with full-file context to drop false positives and write fixes, ships as CLI, VS Code extension, GitHub Action, and pre-commit hook
- Baseline fingerprinting for legacy-codebase adoption, custom YAML rules, watch mode, ~4,400 lines, 138 tests

Akatskii: Cognitive AI Layer

Python, FastAPI, Groq, Claude, Gemini, fastembed

routing extracted to PyPI: llm-route · github.com/AlexlaGuardia/Akatskii

- Hybrid regex + LLM router classifies inputs into 5 thought types and routes by cost/complexity, ~80% resolve instantly and free; agentic tool loop runs up to 15 iterations with parallel calls over MCP
- Semantic memory (fastembed/ONNX, hybrid keyword + cosine recall) and a compaction daemon delivering ~3x context compression across sessions

Guardia Content: AI SaaS Platform

Python, FastAPI, Next.js, React, Stripe

Built end-to-end; public launch pending Meta platform approval

- Social-media automation pipeline: upload → AI styling (SDXL) → caption (Llama 3.3 70B) → quality gate → schedule → multi-platform post; full Stripe billing, Meta OAuth2, and custom domains via Cloudflare-for-SaaS
- Multi-service backend under PM2 with client dashboard, admin HQ, and public per-client microsites

OPEN SOURCE & WRITING

Open source: merged fix into PostHog's LLM-observability product (trace view dropping spans on long-running traces; their AI-observability team adopted the branch and extended it), two merged PRs into PrefectHQ/fastmcp (25.6k★), and an open resource-leak fix to hatchet-dev/hatchet (YC) · **Distribution:** 8 packages published to PyPI and a VS Code Marketplace extension · **Writing:** 8+ technical articles on building MCP servers & AI-agent infrastructure (Dev.to)

ALSO SHIPPED

Stampwerk: live AI freelancer SaaS (stampwerk.com) · **guardia-mcp:** 95+ tools, frame-filtered (50–90% fewer tokens), OAuth 2.1, behind 3 AI interfaces · **MCP Server Suite:** 262 tools across 4 PyPI servers · **Supra:** from-scratch 3D FPS engine in Rust (85K lines, 8 crates) · **Scionbee:** greenhouse WPS-compliance PWA (Next.js, Prisma) · **Paradise:** 4-agent multi-strategy trading system · **mcpcat** & **llm-route:** open-source dev tools (PyPI)